


```
#insert
images_bb[i,:,:] = img

###

#End of my code here

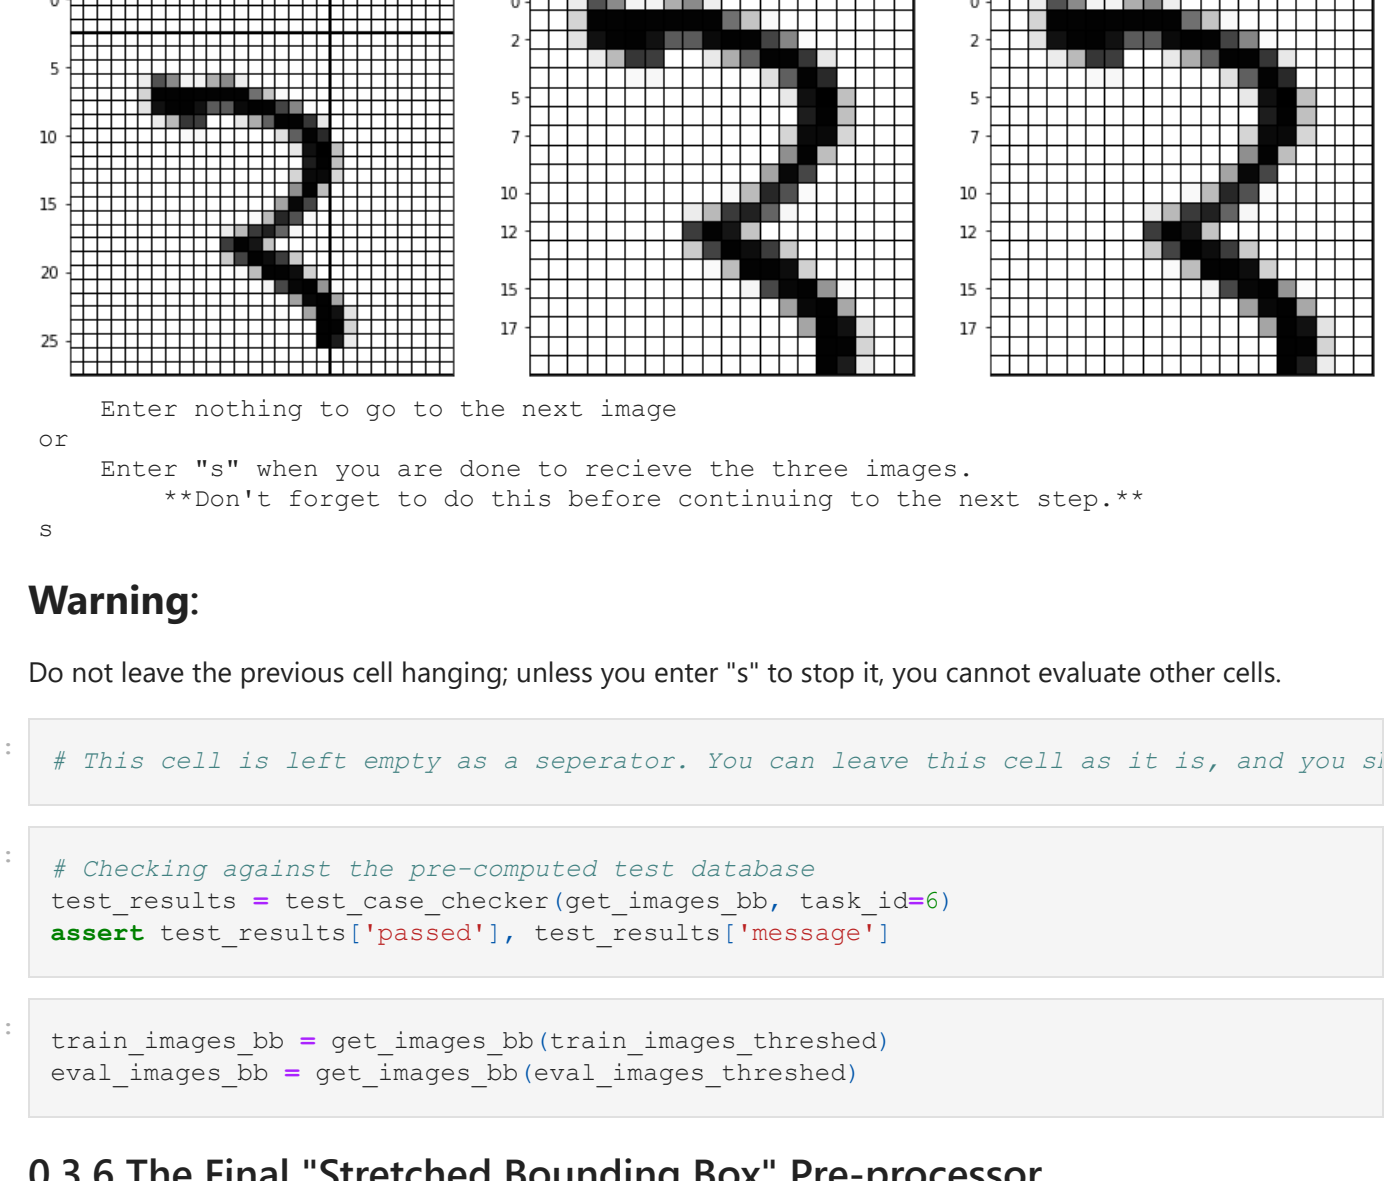
if len(images.shape)==2:
    # In case a 2d image was given as input, we'll get rid of the dummy dimension
    return images_bb[0]
else:
    # Otherwise, we'll just work with what's given
    return images_bb
```

In [36]: # This cell is left empty as a separator. You can leave this cell as it is, and you s

```
In [37]: (orig_image, ref_image, test_im, success_bb) = show_test_cases(get_images_bb, task_id=
```

assert success_bb

The reference and solution images are the same to a T! Well done on this test case.



Enter nothing to go to the next image

or Enter "a" when you are done to receive the three images.

Don't clikc to do this before continuing to the next step.

8

Warning:

Do not leave the previous cell hanging; unless you enter "s" to stop it, you cannot evaluate other cells.

In [38]: # This cell is left empty as a separator. You can leave this cell as it is, and you s

```
In [39]: # Checking against the pre-computed test database
test_results = test_case_checker(get_images_bb, task_id=6)
assert test_results['passed'], test_results['message']
```

In [40]: train_images_bb = get_images_bb(train_images_threshed)
eval_images_bb = get_images_bb(eval_images_threshed)

0.3.6 The Final "Stretched Bounding Box" Pre-processor

Task 7

Similarly, write the function `get_images_sbb` that applies the "Stretched Bounding Box" pre-processing

step and takes following arguments:

- `images`: A numpy array with the shape `(N,height,width)`, where
 - `N` is the number of samples and could be anything.
 - `height` is each individual image's height in pixels (i.e., number of rows in each image).
 - and `width` is each individual image's width in pixels (i.e., number of columns in each image).

Do not assume anything about `images`'s dtype or number of samples.

- `bb_size`: A scalar with the default value of 20, and represents the desired bounding box size.

and returns the following:

- `images_sbb`: A numpy array with the shape `(N,bb_size,bb_size)`, and the same dtype and the range of values as `images`.

The `get_images_sbb` should find a tight-canvas of the ink area in each input image, and stretch it out to fill the full height and width of the output bounding-box. Please see the visual example in the **Assignment Summary** section; the right image should supposedly be the `get_images_sbb` function's output.

We have provided a template function that uses the previous functions and only requires you to fill in the missing parts. It also handles the input shapes in an agnostic way.

Hint: Make sure that you use the `skimage.transform.resize` function from the skimage library. Read about it at <https://scikit-image.org/docs/dev/api/skimage.transform.html>

highlight=resize`skimage.transform.resize`. You may need to pay attention to the `preserve_range`

argument.

skimage.transform.resize(image, output_shape, order=None, mode='reflect', cval=0, clip=True, preserve_range=False, anti_aliasing=None, anti_aliasing_sigma=None) **Resize image to match a certain size.**

Performs interpolation to up-size or down-size N-dimensional images. Note that anti-aliasing should be enabled when downsizing images to avoid aliasing artifacts. For downsampling with an integer factor also see `skimage.transform.downscale_local_mean`.

Parameters image: ndarray. Input image.

output_shape iterable. Size of the generated output image (rows, cols, ...[, dim]). If dim is not provided, the number of channels is preserved. In case the number of input channels does not equal the number of output channels a n-dimensional interpolation is applied.

Channels?

Returns resized. ndarray. Resized version of the input.

```
In [41]: from skimage.util import img_as_uint

def get_images_sbb(images, bb_size=20):
    """
    Applies the "Stretched Bounding Box" pre-processing step to images.

    Parameters:
        images (np.array): A numpy array with the shape (N,height,width)

    Returns:
        images_sbb (np.array): A numpy array with the shape (N,bb_size,bb_size)
        and the same dtype and the range of values as images.

    """
```

if len(images.shape)==2:

In case a 2d image was given as input, we'll add a dummy dimension to be co

images = images.reshape(1,*images.shape)

else:

Otherwise, we'll just work with what's given

images = images

is_row_inky = get_is_row_inky(images)

is_col_inky = get_is_col_inky(images)

first_ink_rows = get_first_ink_row_index(is_row_inky)

last_ink_rows = get_last_ink_row_index(is_row_inky)

first_ink_cols = get_first_ink_col_index(is_col_inky)

last_ink_cols = get_last_ink_col_index(is_col_inky)

#Start my code here

images_sbb = []

for i, img in enumerate(images):

#find a tight-canvas of the ink area in each input image:

img = img[first_ink_rows[i]:last_ink_rows[i] + 1, first_ink_cols[i]:last_ink

#and stretch it out to fill the full height and width of the output bounding

img = resize(img, output_shape=(bb_size,bb_size), preserve_range=True)

images_sbb.append(img)

images_sbb = np.array(images_sbb).astype(images.dtype)

#End of my code here

if len(images.shape)==2:

In case a 2d image was given as input, we'll get rid of the dummy dimension

return images_sbb[0]

else:

Otherwise, we'll just work with what's given

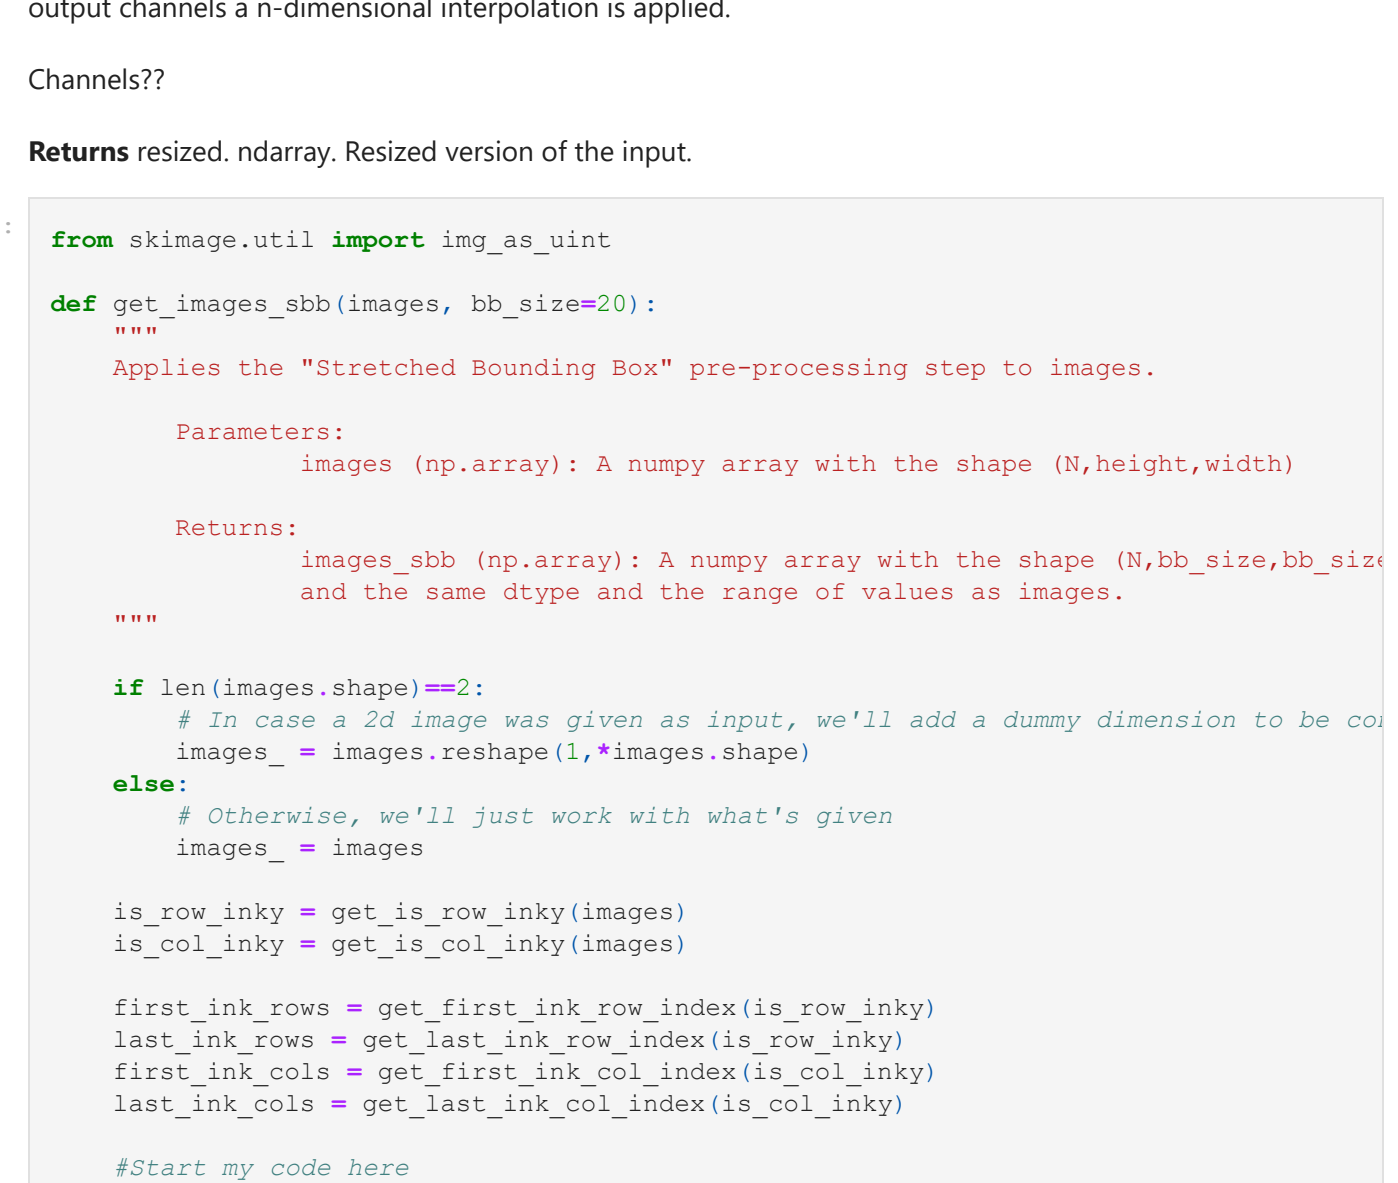
return images_sbb

In [42]: # This cell is left empty as a separator. You can leave this cell as it is, and you s

```
In [43]: (orig_image, ref_image, test_im, success_sbb) = show_test_cases(get_images_sbb, task_
```

assert success_sbb

The reference and solution images are the same to a T! Well done on this test case.



Enter nothing to go to the next image

or Enter "s" when you are done to receive the three images.

Don't forget to do this before continuing to the next step.

8

Warning:

Do not leave the previous cell hanging; unless you enter "s" to stop it, you cannot evaluate other cells.

In [44]: # This cell is left empty as a separator. You can leave this cell as it is, and you s

```
In [45]: # Checking against the pre-computed test database
test_results = test_case_checker(get_images_sbb, task_id=7)
assert test_results['passed'], test_results['message']
```

In [46]: if perform_computation:

print("This is gonna take a while to finish...")

time train_images_sbb = get_images_sbb(train_images_threshed)

time eval_images_sbb = get_images_sbb(eval_images_threshed)

This is gonna take a while to finish...

CPU times: user 39.1 s, sys: 14.1 s, total: 53.2 s

Wall time: 26.7 s

CPU times: user 6.3 s, sys: 2.29 s, total: 8.59 s

Wall time: 4.31 s

1. Naive Bayes Performances

Task 8

Similarly, write the function `train_nb_eval_acc` that trains Naive Bayes models and takes following the

arguments:

- `train_images`: A numpy array with the shape `(N,height,width)`, where
 - `N` is the number of samples and could be anything.
 - `height` is each individual image's height in pixels (i.e., number of rows in each image).
 - and `width` is each individual image's width in pixels (i.e., number of columns in each image).

Do not assume anything about `images`'s dtype or number of samples.

- `train_labels`: A numpy array with the shape `(N,)`, where `N` is the number of samples and has the `int64` dtype.

- `eval_images`: The evaluation images with similar characteristics to `train_images`.

- `eval_labels`: The evaluation labels with similar characteristics to `train_labels`.

- `density_model`: A string that is either 'Gaussian' or 'Bernoulli'. In the former (resp. latter) case, you should train a Naive Bayes with the Gaussian (resp. Bernoulli) density model.

and returns the following:

- `eval_acc`: a floating number scalar between 0 and 1 that represents the accuracy of the trained model on the evaluation data.

We have provided a template function that uses the previous functions and only requires you to fill in the missing parts. It also handles the input shapes in an agnostic way.

Note: You do not need to implement the Naive Bayes classifier from scratch in this assignment; Make sure you use `scikit-learn`'s Naive Bayes module for training and prediction in this task. We have already imported these two functions in the first code cell.

- from sklearn.naive_bayes import GaussianNB, BernoulliNB

```
In [47]: def train_nb_eval_acc(train_images, train_labels, eval_images, eval_labels, density_m
    """
    Trains Naive Bayes models, apply the model, and return an accuracy.

    Parameters:
        train_images (np.array): A numpy array with the shape (N,height,width)
        train_labels (np.array): A numpy array with the shape (N,), where N is
        has the int64 dtype.
        eval_images (np.array): The evaluation images with similar characteri
        eval_labels (np.array): The evaluation labels with similar characteri
        density_model (string): A string that is either 'Gaussian' or 'Bernou
```

Returns:

eval_acc (np.float): a floating number scalar between 0 and 1 that

represents the accuracy of the trained model on the evaluation data.

"""

assert density_model in ('Gaussian', 'Bernoulli')

code start

if density_model == 'Gaussian':

clf = GaussianNB()

else:

clf = BernoulliNB()

why not:

clf = lambda density_model : GaussianNB() if density_model == 'Gaussian' else B

train_images = train_images.reshape(train_images.shape[0], -1)

eval_images = eval_images.reshape(eval_images.shape[0], -1)

Train model X train, Y train

clf.fit(train_images, train_labels)

eval_acc = clf.score(X = eval_images, y = eval_labels)

#code end

return eval_acc

Don't mind the following lines and do not change them

train_nb_eval_acc_gauss = lambda *args, **kwargs: train_nb_eval_acc(*args, density_mod

train_nb_eval_acc_bern = lambda *args, **kwargs: train_nb_eval_acc(*args, density_mod

In [48]: # This cell is left empty as a separator. You can leave this cell as it is, and you s

```
In [49]: # Checking against the pre-computed test database
test_results = test_case_checker(train_nb_eval_acc_gauss, task_id='8 G')
assert test_results['passed'], test_results['message'] # Gaussian Model Test Results

test_results = test_case_checker(train_nb_eval_acc_bern, task_id='8 B')
assert test_results['passed'], test_results['message'] # Bernoulli Model Test Results
```

In [50]: df = None
if perform_computation:

acc_nbq_thr = train_nb_eval_acc(train_images_threshed, train_labels, eval_images_threshed, eval_labels, density_model=

acc_nbb_thr = train_nb_eval_acc(train_images_threshed, train_labels, eval_images_threshed, eval_labels, density_model=

acc_nrb_sbb = train_nb_eval_acc(train_images_sbb, train_labels, eval_images_sbb, eval_labels, density_model='Gaus

acc_nbb_sbb = train_nb_eval_acc(train_images_sbb, train_labels, eval_images_sbb, eval_labels, density_model='Bern

df = pd.DataFrame({'Untouched images', acc_nbq_thr, acc_nbb_thr},

columns = ['Accuracy', 'Gaussian', 'Bernoulli'])

df

Out[50]:

	Accuracy	Gaussian	Bernoulli
0	Untouched images	0.5491	0.8430
1	Stretched bounding box	0.8253	0.8098

2. Decision Forests Performances

Task 9

Similarly, write the function `train_tree_eval_acc` that trains Decision Forest models and takes

following the arguments:

- `train_images`: A numpy array with the shape `(N,height,width)`, where
 - `N` is the number of samples and could be anything.
 - `height` is each individual image's height in pixels (i.e., number of rows in each image).
 - and `width` is each individual image's width in pixels (i.e., number of columns in each image).

Do not assume anything about `images`'s dtype or number of samples.

- `train_labels`: A numpy array with the shape `(N,)`, where `N` is the number of samples and has the `int64` dtype.

- `eval_images`: The evaluation images with similar characteristics to `train_images`.

- `eval_labels`: The evaluation labels with similar characteristics to `eval_labels`.

- `tree_num`: An integer number representing the number of trees in the decision forest.

- `tree_depth`: An integer number representing the maximum tree depth in the decision forest.

- `random_state`: An integer with a default value of 12345 that should be passed to the scikit-learn's classifier constructor for reproducibility and auto-grading (**Do not assume** that it is always 12345).

and returns the following:

- `eval_acc`: A floating number scalar between 0 and 1 that represents the accuracy of the trained model on the evaluation data.

We have provided a template function that uses the previous functions and only requires you to fill in the missing parts. It also handles the input shapes in an agnostic way.

Note: You do not need to implement the Random Forest classifier from scratch in this assignment; Make sure you use `scikit-learn`'s Random Forest module for training and prediction in this task. We have already imported this function in the first code cell:

- from sklearn.ensemble import RandomForestClassifier

- You may need to set 'shuffle = True' due to a known sklearn issue.

clf = RandomForestClassifier(max_depth=2, random_state=0) clf.fit(X, y)

Parameters (n_estimators=100 max_depth=None random_state=None

Methods

- apply(X): Apply trees in the forest to X, return leaf indices.

- fit(X, y): Build a forest of trees from the training set (X, y). Parameters: X "(n_samples, n_features)" y "array-like of shape (n_samples,)"

- score(X, y): Return the mean accuracy on the given test data and labels.

- `random_state`: An integer with a default value of 12345 that should be passed to the scikit-learn's classifier constructor for reproducibility and auto-grading (**Do not assume** that it is always 12345).

- You may need to set 'shuffle = True' due to a known sklearn issue.

```
In [51]: def train_tree_eval_acc(train_images, train_labels, eval_images, eval_labels, tree_num
    """
    Trains Naive Bayes models, apply the model, and return an accuracy.

    Parameters:
        train_images (np.array): A numpy array with the shape (N,height,width)
        train_labels (np.array): A numpy array with the shape (N,), where N is
        has the int64 dtype.
        eval_images (np.array): The evaluation images with similar characteri
        eval_labels (np.array): The evaluation labels with similar characteri
        tree_num (int): An integer number representing the number of trees in
        tree_depth (int): An integer number representing the maximum tree dept
        random_state (int): An integer with a default value of 12345 that sho
        the scikit-learn's classifier constructor for reproducibility and auto
```

Returns:

eval_acc (np.float): a floating number scalar between 0 and 1 that

represents the accuracy of the trained model on the evaluation data.

"""

tree_num = int(tree_num)

tree_depth = int(tree_depth)

random_state = int(random_state)

student code beginning

train_images = train_images.reshape(train_images.shape[0], -1)

eval_images = eval_images.reshape(eval_images.shape[0], -1)

clf = RandomForestClassifier(max_depth=tree_depth, n_estimators=tree_num, random_

clf.fit(X=train_images, y=train_labels)

eval_acc = clf.score(X = eval_images, y=eval_labels)

#student code end

return eval_acc

In [52]: # This cell is left empty as a separator. You can leave this cell as it is, and you s

```
In [53]: # Checking against the pre-computed test database
test_results = test_case_checker(train_tree_eval_acc, task_id=9)
assert test_results['passed'], test_results['message']
```

2.1 Accuracy on the Untouched Images

```
In [54]: df = None
if perform_computation:
```

tree_nums = [10, 20, 30]

tree_depths = [4, 8, 16]

train_images = train_images_threshed

eval_images = eval_images_threshed

acc_arr_uint = np.zeros((len(tree_nums), len(tree_depths)))

for row, tree_num in enumerate(tree_nums):

for col, tree_depth in enumerate(tree_depths):

acc_arr_uint[row, col] = train_tree_eval_acc(train_images, train_labels, ev

tree_num=tree_num, tree_depth=

df = pd.DataFrame([(f'#trees = {tree_num}', *tuple(acc_arr_uint[row])) for row, tree

columns = ['Accuracy'] * [f'depth={tree_depth}' for col, tree_de

print('Untouched Images:')

df

Out[54]:

	Accuracy	depth=4	depth=8	depth=16
0	#trees = 10	0.7496	0.8923	0.9489
1	#trees = 20	0.7707	0.9127	0.9585
2	#trees = 30	0.7883	0.9169	0.9630

2.2 Accuracy on the "Stretched Bounding Box" Images

```
In [55]: df = None
if perform_computation:
```

tree_nums = [10, 20, 30]

tree_depths = [4, 8, 16]

train_images = train_images_sbb

eval_images = eval_images_sbb

acc_arr_sbb = np.zeros((len(tree_nums), len(tree_depths)))

for row, tree_num in enumerate(tree_nums):

for col, tree_depth in enumerate(tree_depths):

acc_arr_sbb[row, col] = train_tree_eval_acc(train_images, train_labels, ev

tree_num=tree_num, tree_depth=

df = pd.DataFrame([(f'#trees = {tree_num}', *tuple(acc_arr_sbb[row])) for row, tree

columns = ['Accuracy'] * [f'depth = {tree_depth}' for col, tree_d

print('Stretched Bounding Box Images:')

df

Out[55]:

	Accuracy	depth = 4	depth = 8	depth = 16
0	#trees = 10	0.7419	0.9043	0.9527
1	#trees = 20	0.7715	0.9162	0.9639
2	#trees = 30	0.7879	0.9248	0.9671

2.3 Accuracy on the "Bounding Box" Images

```
In [56]: df = None
if perform_computation:
```

tree_nums = [10, 20, 30]

tree_depths = [4, 8, 16]

train_images = train_images_bb

eval_images = eval_images_bb

acc_arr_bb = np.zeros((len(tree_nums), len(tree_depths)))

for row, tree_num